

WebAssembly

Wasm On the Edge

Prepared for syd<video> meet-up

Kev Staunton-Lambert

pyrmontbrewery.com


The logo consists of a solid blue square. Inside the square, the letters 'WA' are written in a large, white, bold, sans-serif font. The 'W' and 'A' are closely spaced.

WA

Okay, **what** is Wasm...

Cross platform binary standard to run **pre-compiled code**
from any **web browser*** on any **target platform***

(* major browsers on the platforms they support:
pc/mac/linux)



what else?

... run **pre-compiled** code **server-side** via **Node.js** too!

```
node --expose-wasm
```



and rather excitingly edge compute...

... **sandboxed** user code running on the **CDN edge**!

e.g. manipulate files on-the-fly (not on the origin)



why that's a good thing...

Compiled code means **optimised** code
(typical performance is more than 4x faster than
interpreted JavaScript)

Common op-codes now shared with **hardware architecture**
(improved h/w acceleration options, op-codes limited to
sandbox)

Tiny code, fast to download, small storage/RAM footprint

Uses standard **LLVM** toolchain, so **C / C++, Rust, Go...**
(works well with compiled code but less so Java which relies
on VM to optimise)

why that's a good thing... edge compute

Fastly CDN lets us run WebAssembly on the edge

Apple Low-Latency is a great use case where:

- Playlist can be generated away from the origin
- File parts could be generated rather than uploaded - virtual files chunked on the fly

We can also do other things like modify TS headers on-the-fly

why that's a good thing... edge compute

(example) HLS uses MPEG TS files

TS files are well structured (188 byte packets)

PES PTS/DTS and PCR headers are easily offset on the CDN edge, leaving the origin file untouched

ISOBMFF boxes also easily modifiable

(can be used to remove/smooth out player discontinuities)

how it works: wasm assembler language

Wasm is a binary specification, but includes also a human readable text assembly format

(example) module with a function **getKev** that adds the numbers **2** and **4**

```
kev.wast
```

```
(module
```

```
  (func (export "getKev") (result i32)
```

```
    i32.const 2
```

```
    i32.const 4
```

```
    i32.add
```


how it works: assemble a binary (WebAssembly Binary Toolkit)

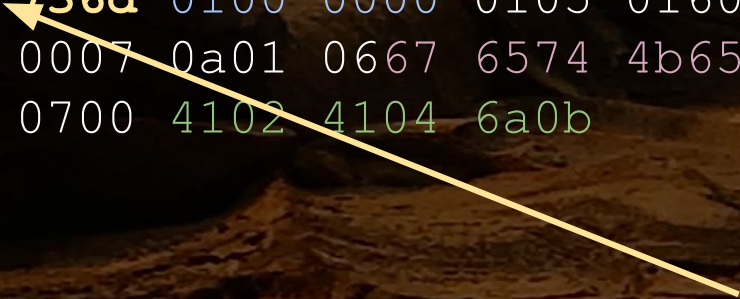
Assemble Wasm text to binary using “Wabbit” tools

```
wat2wasm kev.wast > kev.wasm
```

Creates a small 42 byte binary file called **kev.wasm**

```
0061 736d 0100 0000 0105 0160 0001 7f03  .asm.....`.....
0201 0007 0a01 0667 6574 4b65 7600 000a  ....getKev...
0901 0700 4102 4104 6a0b                ....A.A.j.
```

how it works: what's inside that wasm binary

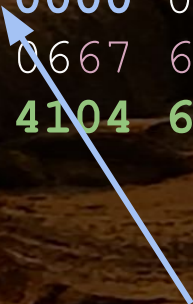


```
0061 736d 0100 0000 0105 0160 0001 7f03 .asm.....`.....
0201 0007 0a01 0667 6574 4b65 7600 000a .....getKev...
0901 0700 4102 4104 6a0b .....A.A.j.
```

Always starts with wasm magic value **0061 736d** (= `.asm`)

how it works: what's inside that wasm binary


0061	736d	0100	0000	0105	0160	0001	7f03	.asm.....`.....
0201	0007	0a01	0667	6574	4b65	7600	000agetKev...
0901	0700	4102	4104	6a0b			A.A.j.



Then wasm version id **0100 0000** (= wasm version 1)

how it works: what's inside that wasm binary


```
0061 736d 0100 0000 0105 0160 0001 7f03  .asm.....`.....
0201 0007 0a01 0667 6574 4b65 7600 000a  ....getKev...
0901 0700 4102 4104 6a0b                ....A.A.j.
```



Then modules: types, exports, such at the function name
Globals,
Tables,
memories,
and

how it works: what's inside that wasm binary

0061	736d	0100	0000	0105	0160	0001	7f03	.asm.....`.....
0201	0007	0a01	0667	6574	4b65	7600	000agetKev...
0901	0700	4102	4104	6a0b			A.A.j.



... **functions** the *opcode instructions section* inside the
getKev exported function:

how it works: those op-codes **4102** **4104** **6a0b**

0x41 is op-code for **push** constant 32-bit integer onto the stack,
we call that twice... 02 + 04

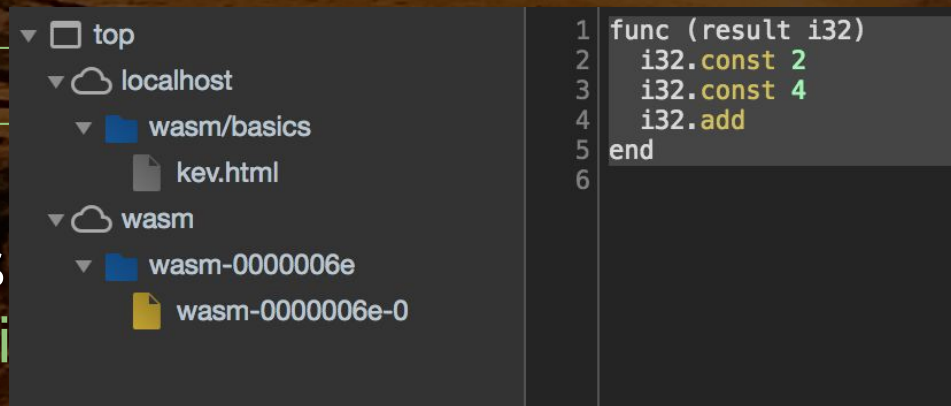
0x**41** **02** -> **i32.const** [**2**]

0x**41** **04** -> **i32.const** [**4**]

0x6a is op-code for **add** from s

0x**6a** -> **i32.add** [i32 i

0x**0b** -> **end**



how it works: more sandboxed op-codes

Many other arithmetic operations and these usual suspects around control:

loop: a block with a label at the beginning which may be used to form loops

if: the beginning of an if construct with an implicit then block

else: marks the else block of an if

br: branch (aka goto) to a given label in an enclosing construct

br_if: conditionally branch to a given label in an enclosing construct

how it works: running wasm from JavaScript

```
<!DOCTYPE html>
<html>
  <body>
    <script>
      WebAssembly.instantiateStreaming(
        fetch('key.wasm')) .then(
          obj => {
            console.log(
              obj.instance.exports.getKev() );
          } );
    </script>
  </body>
</html>
```

Console log shows 6

```
1 <!DOCTYPE html>
2 <html lang="en-AU">
3   <body>
4     <script>
5       WebAssembly.instantiateStreaming(fetch('key.wasm')).then(
6         obj => { obj = {instance: Instance, module: Module}
7         console.log(obj.instance.exports.getKev());
8       });
9     </script>
10  </body>
11 </html>
12
13
```

Object

- ▼ instance: Instance
- ▼ exports: Object
 - ▼ getKev: f 0()
 - arguments: null
 - caller: null
 - length: 0
 - name: "0"
 - ▶ __proto__: f ()
 - ▶ [[Scopes]]: Scopes[0]
 - ▶ __proto__: WebAssembly.Instance
- ▼ module: Module
 - ▶ __proto__: WebAssembly.Module
 - ▶ __proto__: Object

how safe is it? All seems a bit dangerous...

It's not so bad, many implementations so expect some holes, but also quicker patching with community resolution

All memory/table marshalling well defined and type checked

Modules runs in a sandbox and will simply throw a JavaScript exception if something unexpected/naughty happens

Also utilises clang CFI (Control Flow Integrity) checking

<https://clang.llvm.org/docs/ControlFlowIntegrity.html>

what doesn't work so well *yet*

Debugging symbols are currently missing

You *can* debug Wasm e.g. with WebtKit inspector, but this is not stepping the original sources (gdb) which can make things tricky, printf remains your friend

Shared objects (libraries)

You can't link libraries, everything needs to be compiled into same global module right now

Support in SmartTVs, Roku, tvOS, Chromecast etc

Someday maybe :-\ Embedded systems like these would

tools: emscripten: C/C++ -> Wasm

Emscripten modifies LLVM (clang) output into Wasm

```
brew install emscripten
```

Works by overriding LLVM in ~/.emscripten

```
LLVM_ROOT =
```

```
'/usr/local/opt/emscripten/libexec/llvm/bin'
```

```
kev.c
```

```
#include <stdio.h>
```

```
int main(int argc, char ** argv) {
```

```
    printf("G'day, Switch!\n");
```

```
}
```

tools: webassembly.studio: Rust -> Wasm

webassembly.studio is web based IDE tool

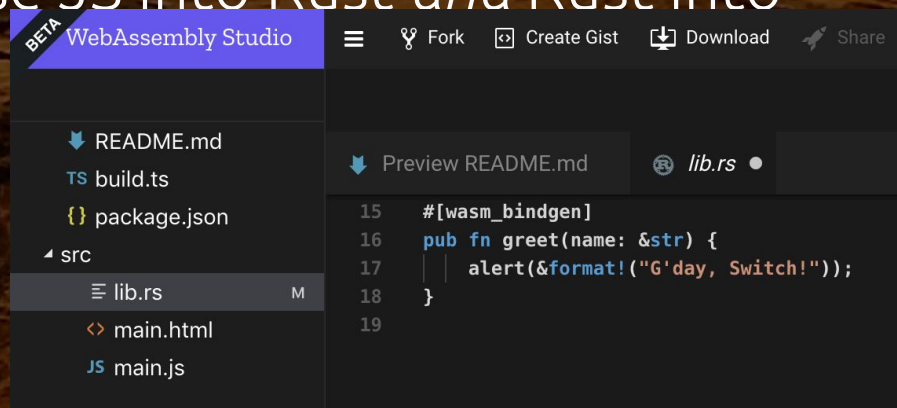
<https://webassembly.studio/>

It uses **wasm_bindgen** to generate the language bindings (which works both ways, expose JS into Rust *and* Rust into JS)

```
kev.rs
```

```
#[wasm_bindgen]
```

```
pub fn greet(name: &str) {  
    alert(&format!("G'day, Switch!"));  
}
```



tools: wasm-bindgen (JS to Rust bridge)

Installs using Rust's Cargo/crates.io package manager

```
curl https://sh.rustup.rs -sSf | sh
```

```
source $HOME/.cargo/env
```

```
rustup -v install nightly
```

```
rustup target add wasm32-unknown-unknown --toolchain nightly
```

```
cargo +nightly install wasm-bindgen-cli
```

Make new project

```
cargo +nightly new kev_rust --lib
```

Then create Cargo.toml to pull Wasm package

tools: wasm-bindgen (Rust to JS bridge)

```
kev_rust/Cargo.toml
[package]
name = "kev_rust"
version = "0.1.0"
authors = ["Kev <keveyski@gmail.com>"]
```

```
[dependencies]
wasm-bindgen = "0.2"
```

Build Wasm

```
cargo +nightly build --target wasm32-unknown-unknown
wasm-bindgen target/wasm32-unknown-unknown/debug/kev_rust.wasm
```


tools: some other Wasm editors

VSCode has a WebAssembly extension

<https://webassembly.studio/>

vi


<https://github.com/rhysd/vim-wasm>

WebKit inspector does pretty good job too (wasm/wast conversion etc)

wasm example: live interactive gaming

Unity3D WebGL

<https://webassembly.org/demo/Tanks/>



wasm example: CDN Edge compute

Modifying MPEG transport files on the edge
Manifest manipulation (ad insertion grep)

future ideas (disclaimer: Kev's thoughts)

Skipping validation step by signing the Wasm binary?

There is a setup cost to running Wasm where it needs to check things are good each time it loads a binary, what if we can confirm we did that already with some sort of **code sign**? Code sign or registry check etc of course needs to be faster than validation itself without becoming an attack vector ;-)

Pre-known metadata could lead to less or **zero JavaScript**

Thanks!

Slides are here:

<https://goo.gl/2ahsEY>

Apple Low-Latency

<https://tinyurl.com/yyr2rz8m>

AV1 <https://goo.gl/pGnNgJ>

wasm@pyrmontbrewery.com.au

